

SUBMITTED IN PARTIAL FULFILLMENT OF REQUIREMENTS FOR THE DEGREE
OF MASTER OF TECHNOLOGY

AMAN MADAAN



OCTOBER, 2014

Distant Supervision for Numerical Relation Extraction

October 24, 2014

Contents

1	Introduction	4
1.1	Terminology	4
2	Problem	4
3	Snowballing	5
4	Distant Supervision: Snowball scaled up	5
4.1	Introduction	5
4.2	Distant Supervision Assumption	6
4.3	Matching	6
4.4	Training	6
5	Survey	6
5.1	Constructing Biological Knowledge Bases by Extracting Information from Text Sources, Craven and Kumlien 1999, [3]	7
5.2	Distant Supervision for the Web, Mintz et. al 2009 [8]	7
5.3	Learning 5000 relation extractors [5]	7
5.4	Riedel et. al [9]	8
5.5	MultiR [6]	8
5.6	Other Works	9
6	Numerical Relation extraction using MultiR	9
6.1	Relations	9
6.2	Knowledge base	10
7	L’homme propose, et Dieu dispose	10
7.1	Numbers and False positives	11
8	Fighting False Positives with units	11
8.1	Results	12

9	Numbers are weak entities: a case for keywords	13
9.1	Motivation	13
9.2	Numerical Relations are Explicit	13
9.3	Approach	13
9.4	Results	14
10	A summary of the three heuristics	14
11	Rule Based Extraction	15
11.1	Peculiarity of Numerical Relations	15
11.2	Dependencies	15
11.3	Dependency Path	16
11.4	Relation extraction using dependency paths	16
11.5	Example	16
11.6	Results	17
	11.6.1 Precision and Recall	17
	11.6.2 Example Extractions	17
12	Conclusion	17

Abstract

The task of numerical relation extraction poses new, hitherto untackled challenges. The bewildering amount of false positives, units, modifiers, time varying relations are just some of the issues that are non existent for standard relation extraction, but become crucial when numbers are involved.

We discuss molding distant supervision for numerical relation extraction. The standard one-to-one mapping using numbers as the second entity fails. Adding units help in improving the training data to a certain extent. The third heuristic, keyword based pruning further improves precision in the matches. Analysis of the results from the three heuristics lead to a basic rule based extractor, which performs better than any of the heuristics.

1 Introduction

Massive knowledge bases containing information encompassed by the Web in a neat, ready to process, structured form continue to be a part of an IR researcher's reverie.

This is not unexpected; such knowledge bases have potential of revolutionizing the way in which information is searched by users on the web or exchanged by machines among themselves. Clearly, any progress towards a solution will have to deal with intricacies of how facts are expressed in the natural language. It turns out that such intricacies are too many to exhaust. Years of research has went into the aforementioned goal, and still there are plenty of loose screws. This report discusses one of them, numerical relation extraction.

We start with the problem definition, followed by an explanation of Distant supervision. We then present a summary of some of the key works. Subsequently, we discuss casting standard distant supervision for numerical relation extraction and the improvements achieved using Units and keywords. Finally, we discuss a simple rule-based relation extractor.

1.1 Terminology

- **a)Entity** An entity is something that exists in itself, actually or hypothetically.¹
- **b)Relation** A relation specifies a concept which binds two entities. For example, creator(Linux, Linus Torvalds).
- **c)Mention** A piece of text which expresses a relation. For example, the sentence "Linus Torvalds is the creator of Linux."
- **d)Match** A mention expressing a relation R is called a match for R . The criteria for deciding whether the mention m is a match or not can vary. Presence of both the entities is, however, a mandatory condition.
- **e)Extraction** A 3-tuple $R(A, B)$ where A and B are entities related via a relation R .

2 Problem

Train extractors that can harness the Web for numerical relations, where relations are 3-tuples linking an entity to a number. For example,

- (India, **economy**, 1.842 trillion USD)
- (China, **internet users**, 590.56 million)
- (USA, **land area**, 2,959,054 square mile)

¹from <http://en.wikipedia.org/wiki/Entity>

3 Snowballing

Let us motivate the idea by considering the following related problem:

Suppose we want to populate a repository of founders of companies, and all that we know is that Elon Musk is the founder of SpaceX. The problem can be divided into the following two parts, each of them rely on an intuition about how human beings form sentences in general.

- *Given an entity pair, and a corpus of documents, find out all the sentences that **express a relation** between the entity-pair.*

Command line ninjas will quickly think of the following solution:

```
grep -i "entity1" sentences | grep -i "entity2"
```

The intuition behind this perhaps the most obvious solution is that *a sentence that houses both the entities can be expected to express a relation between them*. A quick web search with the query *"entity1" and "entity2"* will show that this intuition has some weight.

- *Find what makes these sentences special*

Sentence structure depends on the relation being expressed. In verbose, if two sentences express the same relation, we can expect presence of some *features* that are similar in both of them. These include POS tags, words around the entities and dependency path between the entities to name a few.

Putting together the intuitions above, we can solve the problem as follows:

1. Collect all the sentences which have SpaceX and Elon Musk in them, extract features from these sentences, favor² those features which repeat.
2. Now use this set of features to extract similar pairs from other sentences. A fancier solution would be to reuse the extractions to learn more features, and continuing the process till the point of diminishing returns.

This seemingly shaky method actually works [1] and is popular by the name of snowballing.

4 Distant Supervision: Snowball scaled up

4.1 Introduction

If the idea of Snowball looks convincing, Distant supervision should follow naturally. Earlier, we considered only one relation for one entity pair. Scale the amount of both of these up and we have distant supervision. The basic setup is as follows:

- (a) **KB:** A knowledge base consisting of facts. The facts are 3-tuples; the entities and the corresponding relation. For example:

²By giving them more weight, pruning features that are rare etc.

Entity	Entity	Relation
Donald Knuth	Wisconsin	Born In
Srinivasa Ramanujan	Erode	Born In
Alan Turing	London	Born In
Alon Musk	SpaceX	Founder Of

has 4 different facts

- (b) **Corpus** The repository of text where we expect to find the sentences that express facts that we know. We need another repository, called the test set, where we will run our extractor to obtain new facts. These two can be the same.

4.2 Distant Supervision Assumption

Every sentence that has an entity pair (e_1, e_2) expresses the relation which exists between (e_1, e_2) .

4.3 Matching

We next need to align our knowledge base with the corpus. This process is also called matching.

Data: Corpus C, Knowledge Base KB

Result: Training data, D, A set of matches

Break C into a set of sentences, S;

for each sentence s in S **do**

 let E = all entity pairs in s;

for each entity pair (e_1, e_2) in E **do**

if \exists relation r in KB with $r(e_1, e_2)$ **then**

 add s to D with label r

end

end

end

Algorithm 1: Distant Supervision: Matching

4.4 Training

Recall that obtaining the sentences that express a relation gives us training data, which we want to use to learn relation extractors, our goal. There are several ways to achieve this, starting from the naive ways of training sentence level classifier extending to fancier graphical model based learning. There are broadly two different ways of used in the literature for training extractors using the distant supervision data. Earlier work on distant supervision used the training data to train classical classifiers, like naive bayes, LR and so on. Recent work has shifted towards modeling extraction using graphical models.

5 Survey

The first distant supervision paper came out in 1999 [3]. The process has remained more or less the same since the first paper, with works stemming out

of either relaxations of different assumptions made by [3] or by using a smarter matching algorithm [10].

5.1 Constructing Biological Knowledge Bases by Extracting Information from Text Sources, Craven and Kumlien 1999, [3]

This was the first work to use distant supervision for creating a repository of biological facts. They targeted 5 different relations between Proteins, Tissues, Cell-Types, Diseases, Pharmacologic-Agents and Subcellular-structure. A naive bayes based simple relation extractor is first described. The extractor works in 2 steps: A classifier trained on hand labeled data first labels whether a sentence *can* express a particular relation. If yes, the sentence is searched for a pair of one of the 5 types depending on the label and the fact is added to the database. The authors note that it took an expert 35 hours to hand label the corpus. This forms the basis of motivating distant supervision based methods. The paper also identifies many areas of improvements, like a pair of entities taking up multiple relations, on which subsequent work has been built up.

The authors obtain an improvement of around 9 precision points with distant supervision. It is possible that the following points contributed to the improvement in scores:

- Constrained entity set. Proteins and Tissues won't just appear together without any possible relation.
- The corpus used was very well aligned with the knowledge base.

5.2 Distant Supervision for the Web, Mintz et. al 2009 [8]

This work revived the interest of the community around the problem of distant supervision. The major contributions of this paper can be listed as follows:

- **Distant Supervision for the Web** Craven and Kumlien had a limited knowledge base and a limited corpus to align it with. As the web exploded, the knowledge bases that became available were of the magnitude of Free-Base, and the corpus that can be aligned with it was the web. Mintz et. al bought this fact to the spotlight and sparked a series of works.
- **Features** They designed a set of rich features that are used by researchers to date. These features include dependency paths, POS tag sequences, word sequences to name a few.

5.3 Learning 5000 relation extractors [5]

The number of relations typically used in the works were limited to the range of 30-40. This work targeted 5000 relations, which is a big jump. To fight sparsity that will arise in the training data due to a large number of classes, they also train a top level classifier which decides which extractor should be used for a particular document. Not all the extractors are employed for a given document.

5.4 Riedel et. al [9]

Distant supervision assumption does not really holds when the knowledge base is not well aligned with the corpus. Another way of saying this would be that the corpus can be expected to consist of sentences that can host entity pairs in a wide range of contexts, and the sentences may have nothing to do with the relations in which the entity pairs appear. They take the example of the relation “nationality”. A wide range of popular entities that are born in the country will be mentioned in the sentences that may have nothing to do with the fact that they were born in that country. In such cases, it can be argued that the training data generated will be extremely noisy and will lead to poor extraction performance. Riedel et. al relax the distant supervision assumption.

5.5 MultiR [6]

Some of the entity pairs can occur in multiple relationships. For example, an athlete usually *plays for* the country they are *born in*. The *CEO of* a company also *works for* it. The techniques discussed till now are focused on extracting only one of the relations for an entity pair. MultiR relaxes one relation per entity-pair assumption. (Hence the name, multi-relation). MultiR achieves this by modeling both sentence level and corpus level extraction decision. The graphical model is as shown in figure 1. (taken from [6])

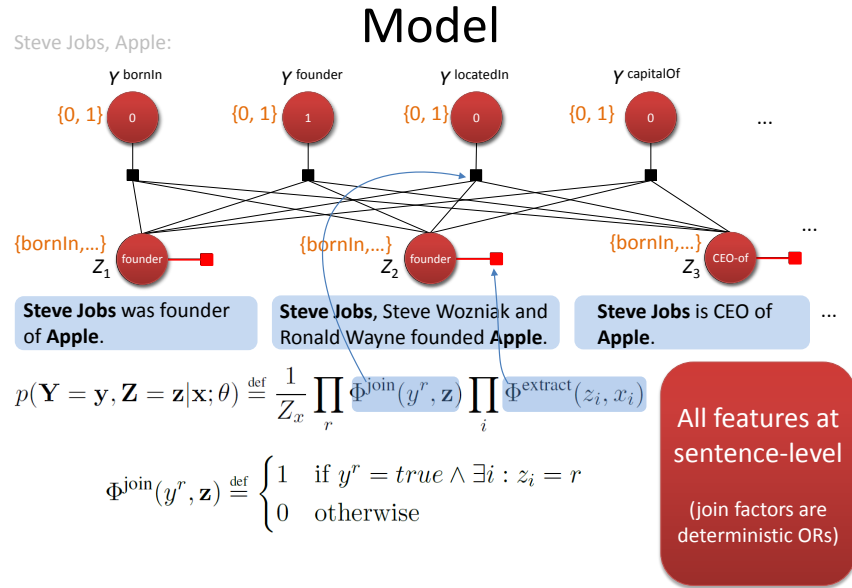


Figure 1: MultiR Graphical Model

There are two different types of nodes in the model:

- **a) Sentence Level extractions, \mathbf{Z} :** The cardinality of these variables is equal to the number of targeted relations + 1 for “no attachment”,

which means that the sentence expresses no relation. This is important for relaxation of the DS assumption.

- **b) Corpus level extraction, \mathbf{Y} :** The Y s are the binary variables, one for each relation. Y_i takes the value of 0 or 1 depending on whether the i th relation is expressed in the corpus or not.

There are 2 factors:

- **a) Factors local to \mathbf{Z}** These factors capture the likelihood of a sentence belonging to one of the relations. The usual technique of representing factor tables in log-linear forms is used. The features are as defined by [8].
- **b) Factors between \mathbf{Y} and \mathbf{Z}** These factors ensure that the probability mass is divided only amongst the possible extractions. As shown in the figure, the factor takes a value of 1 only if the corpus level bit is 1 i.e. that relation exists at corpus level **and** at least one of the sentences also expresses it.

5.6 Other Works

As mentioned, many possibilities in the distant supervision framework have already been explored. Zhou et. al [12] present an alternative model, MIML, that relaxes both distant supervision and one relation per entity pair assumption (steve jobs was both CEO and founder of Apple) much like MultiR. The training in MIML is logistic regression based but is quite time consuming. MultiR employs a perceptron style training algorithm and thus runs much faster. Min et. al [7] argue that since the knowledge bases are incomplete, the extent of false negatives is quite large. Koch et. al [10] add type awareness and coreference resolution during matching and extraction. Takamatsu et. al [11] try to improve quality of training data by learning patterns which cannot possibly express a relation.

6 Numerical Relation extraction using MultiR

Numerical relations are much like the usual entity-entity relations, just more problematic. What we call a *numerical relation* will colloquially be called an attribute. For example, Camel and 215cm are related via the relation “Average Height”.

The problem of numerical relation extraction can be easily cast into the distant supervision framework. Recall from section 4 that we need a corpus and a knowledge base to generate the training data for distant supervision. We also need to fix on the set of relations that will be targeted.

6.1 Relations

The following 10 relations were chosen for our experiments:

Relation Name	Relation Code
Land area (sq. km)	AG.LND.TOTL.K2
Foreign direct investment, net (current US\$)	BN.KLT.DINV.CD
Goods exports (current US\$)	BX.GSR.MRCH.CD
Electricity production (kWh)	EG.ELC.PROD.KH
CO2 emissions (kt)	EN.ATM.CO2E.KT
Pump price for diesel fuel (US\$ per liter)	EP.PMP.DESL.CD
Inflation, consumer prices (annual %)	FP.CPI.TOTL.ZG
Internet users (per 100 people)	IT.NET.USER.P2
GDP (current US\$)	NY.GDP.MKTP.CD
Life expectancy at birth, total (years)	SP.DYN.LE00.IN
Population (Total)	SP.POP.TOTL

Table 1: Selected Relations

The relations were selected keeping in mind the availability of sentences which potentially express these facts and to have some diversity in units.

6.2 Knowledge base

Our knowledge base was derived from data.worldbank.org. It has 4371979 numerical facts about 249 countries ranging over 1281 different attributes. Table 2 shows a sample.

Country id	Number	Relation
/m/04g5k	3126000130	EG.ELC.PROD.KH
/m/02k8k	1969.179	EN.ATM.CO2E.KT
/m/06nnj	332315	SP.POP.TOTL
/m/019rg5	55.020073	SP.DYN.LE00.IN
/m/05sb1	19974.148	EN.ATM.CO2E.KT
/m/05v8c	10000000000	EG.ELC.PROD.KH
/m/03spz	7639000100	EG.ELC.PROD.KH
/m/06vbd	44249.688	EN.ATM.CO2E.KT
/m/0d060g	51.3	IT.NET.USER.P2
/m/05qkp	62.298927	SP.DYN.LE00.IN

Table 2: A Sample of our KB

7 L’homme propose, et Dieu dispose

Our life would have been simpler if MultiR would do the desired; extract the numerical relations from the corpus provided information about some of the numerical relations. However, as table 3 shows, the quality of training data is so poor that not much can be expected from the extractor obtained.

Relation	Total Matches	Sampled Matches	True Matches	Precision(%)
Land Area	1884	15	1	6.7
Foreign Direct Investment	0	0	0	0
Goods Export	0	0	0	0
Electricity Production	381	10	0	0
CO ₂ Emission	0	0	0	0
Diesel Prices	8491	15	0	0
Inflation(%)	8689	15	0	0
Internet Users(%)	182319	40	0	0
GDP(\$)	0	0	0	0
Life Expectancy	267	10	0	0
Total Population	0	0	0	0

Table 3: Vanilla Match

It turns out that vanilla distant supervision leads to an **unprecedented** amount of false positives, which are the root cause of everything that went wrong in the distant supervision pipeline.

7.1 Numbers and False positives

Why would numbers lead to false positives in the first place? The problem stems from the fact that numbers have no identity of their own; they represent count of some real entity or phenomenon.

- **Numbers can appear in many more contexts with an entity** The number of ways in which any two entities can appear together in a sentence is far less than the number of ways in which a number and a quantity can appear together. For example, Consider the entity pair “Bill Gates” and “Microsoft” and the entity-number pair “Bill Gates” and “3” (say). While former will usually co-occur in finite contexts (Founder, CEO, Evangelist etc.), the latter may co-occur anywhere Bill Gates happen to be around something which is 3, the number of cars, billion dollars donated, number of units headed, position in the company, number of business units shutdown by Microsoft and so on.
- The situation is worse for smaller whole numbers, which are more frequent. This intuitively makes sense as we are more often see 2,3 or 11 than 111212233 or 11.42143.

- **The match mines**

During the initial phases of our experiments, we stumbled upon the *match mines* These were basically huge tables, world cup scores of all the matches played and so on. A couple of such sentences were responsible for 21% of the matches! It is easy (and very important) to get rid of such sentences. For subsequent runs, we first sort the sentences by length and then remove top 1000 of them.

8 Fighting False Positives with units

Analyzing results of plain matching made it clear that units will help in improving both precision (by eliminating matches where the unit is not present)

and recall (increasing matches by canonicalization of numbers and conversion to SI units). We found that though units helped in drastically cutting down the number of false positives (match mines were completely eliminated), and helped recall (lots of good matches for Land area and Population), the number of false positives was still a trouble. The number of false positives was typically high for cases where the unit was percentage, since it is again a very generic unit. For other relations too, the number of false positives was very large. The large number of false positives, apart from degrading quality of the model, make evaluating the quality of matcher very difficult.

As the next step, we integrated Prof. Sunita’s unit extractor in the distant supervision pipeline. Concretely, we fed the corpus to the unit extractor to get information of the following form (sample):

Sentence ID	Numbers with units
2000064	172010.0::64:72;500.0::157:160;
2000077	75.0::united states dollar:17:19;2010.0::55:67;29.0::69:70;
2000080	50.0::47:49;
2000112	172009.0::56:64;
2000113	0.10000000149011612::58:61;
2000120	3.15564E7::second:113:115;1.0::131:132;

Where the format of a single match is:

sentence id [TAB] Number, unit, startOffset, endOffset;Number2,unit,startOff,endOff

The knowledge base was also converted to SI units. We see lots of true positives for attributes that take value in the higher range. The number of false positives has also drastically decreased. Units have been particularly helpful in dealing with the so called “match mines”, sentences which have lots of country names and numbers, like score tallies.

8.1 Results

As table 4 shows, units help, but not as much as we would have expected.

Relation	Total Matches	Sampled Matches	True Matches	Precision(%)
Land Area	98	40	32	80
Foreign Direct Investment	791	40	1	2.5
Goods Export	816	40	3	7.5
Electricity Production	19	19	0	0
CO ₂ Emission	196	40	2	5
Diesel Prices	2	2	2	100
Inflation(%)	27598	40	0	0
Internet Users(%)	24639	40	0	0
GDP(\$)	1790	40	0	0
Life Expectancy	3081	40	0	0
Total Population	5225	40	11	27.5

Table 4: Unit Based

There can be several reasons for this.

- Attributes with Unit % (inflation, internet user %), are affected by stock data that is heavily present in the news corpus.

- Relations FDI, Goods Export, GDP have almost similar values, thus, an (entity, number) pair matches all of the three relations.

The analysis indicates that we might need some relation specific heuristics.

9 Numbers are weak entities: a case for keywords

9.1 Motivation

Matches from unit extraction showed that in some cases, the sentence is labeled as a match for a particular relation has no mention of the relation itself at all. For example, consider: *“In eurozone powerhouse Germany, industrial orders jumped 3.2 percent in June, official data showed Thursday, with foreign demand behind a sharp rebound following a surprise drop in May.”* In this sentence, (Germany, 3.2) was considered as a match pair for the relation *Internet user percent*. Clearly, it has nothing to do with it.

9.2 Numerical Relations are Explicit

A key observation that can be made by going through the sentences that express numerical relations is that one cannot be too poetic while forming a sentence that is supposed to state a numerical fact. This is in stark contrast with sentences expressing relations between entity pairs, wherein the underlying relation might be implicit. If we want to state GDP of a country in a sentence, there is no escape from the words like “GPD” or “gross domestic product” and the likes ³.

If this is indeed true, imposing an additional constraint of a keyword being present in a sentence in addition to the fact being present can help in cutting down the number of false positives. We note that such a pruning is possible only in case of numerical relations. As mentioned earlier, for real world entity pairs, co-incidental matches will be rarer and a constraint on the relation word being present will be too restrictive.

9.3 Approach

Let M_r be the set of matches obtained by standard unit + distance based matching for a relation r . We prune M_r by picking only sentences which contain one of the words in the set $\text{keywords}(r)$. The sets $\text{keywords}(r)$ are manually crafted and are listed in table 5.

³There can be a few exceptions for unit-less quantities like “Dell, 49, lives with his family”

Relation Keywords	(case insensitive)
Internet User %	"Internet"
Land Area	"area", "land", "land area"
Population	"Population"
Diesel	"diesel"
GDP	"Gross domestic", "GDP"
CO2	"Carbon", "Carbon Emission", "CO2"
Inflation	"Inflation", "Price Rise"
FDI	"Foreign", "FDI"
Goods Export	"goods"
Life Expectancy	"life", "life expectancy"
Electricity Production	"Electricity"

Table 5: Keywords for Relations

9.4 Results

The results indicate that keywords help a lot.

Relation	Total Matches	Sampled Matches	True Matches	Precision(%)
Land Area	61	40	30	75
Foreign Direct Investment	8	0	0	0
Goods Export	4	4	0	0
Electricity Production	0	0	0	0
CO ₂ Emission	16	16	2	12.5
Diesel Prices	2	2	2	100
Inflation(%)	3853	40	25	62.5
Internet Users(%)	308	40	7	17.5
GDP(\$)	0	0	0	0
Life Expectancy	99	40	15	37.5
Total Population	607	40	24	60

Table 6: Keyword Based

10 A summary of the three heuristics

Before discussing the results, let us summarize the heuristics:

- **a) Vanilla:** A sentence having a country-number pair $c - n$ is called a match for a relation R if $R(c, n)$ is a fact in the KB.
- **b) Unit Based:** Feed the corpus to a unit extractor, extract normalized numbers along with units. During matching, both the normalized value of a number and the unit are considered.
- **c) Keyword based:** In addition to matching number and unit, also check whether one of the keywords for a relation is present.




Figure 2: A comparison of the three heuristics

11 Rule Based Extraction

11.1 Peculiarity of Numerical Relations

Analyzing a number of sentences expressing numerical relations lead to several insights as already discussed.

- **Keywords** Sentences expressing numerical relations can be expected to be explicit about the relation being expressed. Stated another way, we can expect presence of certain keywords that might help in identifying relations.
- **Modifiers** A large number of false positives stem out of mentions where a change in the numerical attribute is mentioned.

11.2 Dependencies

Dependencies are grammatical relation between two words, governor and dependent. The relation captures the way in which one of the words is affected by the other. For example, consider the sentence: “The red ball was lost” The dependencies are:

- **amod(ball,3,red,2)** “Red” is an adjective for “ball”
- **det(ball,3,The,1)** “the” is a determiner of “ball”
- **nsbjpass(lost,5,ball,3)** “ball is the subject of lost”




Figure 3: Dependency Graph

- **auxpass(lost,5,was,4)** “was is an auxiliary of lost”

Refer [4] for a detailed explanation of all the dependencies with examples.

11.3 Dependency Path

We define a dependency path between two words “A” and “B” as the shortest path between them in the dependency graph. The dependency graph consists of one node for each of the words, and the dependencies are the collapsed typed dependencies as obtained from the stanford dependency parser. [2]. Figure [fig:1] shows dependency graph for “The estimated 2014 population of Zambia is 15,200,000, which ranks 70th in the world.”

11.4 Relation extraction using dependency paths


Intuitively, it makes sense that the entities which are related will have some dependence on each other in the sentence expressing the relation. If this is true, we can expect the following simple rule to work. Let P be the set of words in the path between a Country C and a number N . Then, for C - N to be an extraction for a relation R , the following should hold:

- **P should have** a word w such that $w \in keywords(R)$
- **should not have** any word w such that $w \in modifiers$

Here *modifiers* is a set of modifying words. The set is formed of the words like “Increased, Decreased, Changed”.

11.5 Example


Consider the sentence: “The estimated population for 2014 of the Australian continent is about 36.25 million people” and the Country-Number pair (Australian, 36.25) The dependency graph of the sentence is as shown in figure [pos].



The path between Australia and 36.25 has “people”, so this will be an extraction.

If we modify the sentence to: “The estimated population for 2014 of the Australian continent increased by about 3.25 million people”

The path will additionally have a modifier, “increased”, and thus there won’t be any extraction.



11.6 Results

11.6.1 Precision and Recall

The extractor was applied to 30 sentences expressing 23 different relations.

	Relations Present	Relations not Present (False positives)
Extracted	16	17
Not Extracted	7	N/A


- Precision: 48.4%
- Recall: 69.6%

The precision should increase further on applying unit based pruning.

11.6.2 Example Extractions

Sentence → Extraction
At 3.71 million square miles (9.62 million km2) and with around 318 million people, the US is the world’s 3rd or 4th-largest country by total area and third-largest by population. → POP(US, 318)
The land area [hoffmann2011knowledge] of the contiguous US is 2,959,064 square miles (7,663,941 km2) → AGL(US, 2,959,064), AGL(US, 7,663,941)
With 1,210,193,422 residents reported in the 2011 provisional census, India is the world’s second-most populous country. → POP(India, 1,028,737,436)
According to an official estimate for 1 June 2014, the population of Russia is 143,800,000. → POP(Russia, 1), POP(Russia, 2014), POP(Russia, 143,800,000)

12 Conclusion



References

- [1] Eugene Agichtein and Luis Gravano. “Snowball: Extracting relations from large plain-text collections”. In: *Proceedings of the fifth ACM conference on Digital libraries*. ACM. 2000, pp. 85–94.
- [2] Stanford CoreNLP. *A Suite of Core NLP Tools*.
- [3] Mark Craven, Johan Kumlien, et al. “Constructing biological knowledge bases by extracting information from text sources.” In: *ISMB*. Vol. 1999. 1999, pp. 77–86.
- [4] Marie-Catherine De Marneffe and Christopher D Manning. “Stanford typed dependencies manual”. In: *URL http://nlp.stanford.edu/software/dependencies_manual.pdf* (2008).
- [5] Raphael Hoffmann, Congle Zhang, and Daniel S Weld. “Learning 5000 relational extractors”. In: *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics. 2010, pp. 286–295.
- [6] Raphael Hoffmann et al. “Knowledge-based weak supervision for information extraction of overlapping relations”. In: *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies-Volume 1*. Association for Computational Linguistics. 2011, pp. 541–550.
- [7] Bonan Min et al. “Distant Supervision for Relation Extraction with an Incomplete Knowledge Base.” In: *HLT-NAACL*. 2013, pp. 777–782.
- [8] Mike Mintz et al. “Distant supervision for relation extraction without labeled data”. In: *Proceedings of the Joint Conference of the 47th Annual Meeting of the ACL and the 4th International Joint Conference on Natural Language Processing of the AFNLP: Volume 2-Volume 2*. Association for Computational Linguistics. 2009, pp. 1003–1011.
- [9] Sebastian Riedel, Limin Yao, and Andrew McCallum. “Modeling relations and their mentions without labeled text”. In: *Machine Learning and Knowledge Discovery in Databases*. Springer, 2010, pp. 148–163.
- [10] Mitchell Koch John Gilmer Stephen Soderland and Daniel S Weld. “Type-Aware Distantly Supervised Relation Extraction with Linked Arguments”. In: *Proceedings of EMNLP 2014*.
- [11] Shingo Takamatsu, Issei Sato, and Hiroshi Nakagawa. “Reducing wrong labels in distant supervision for relation extraction”. In: *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics: Long Papers-Volume 1*. Association for Computational Linguistics. 2012, pp. 721–729.
- [12] Zhi-Hua Zhou et al. “Multi-instance multi-label learning”. In: *Artificial Intelligence* 176.1 (2012), pp. 2291–2320.