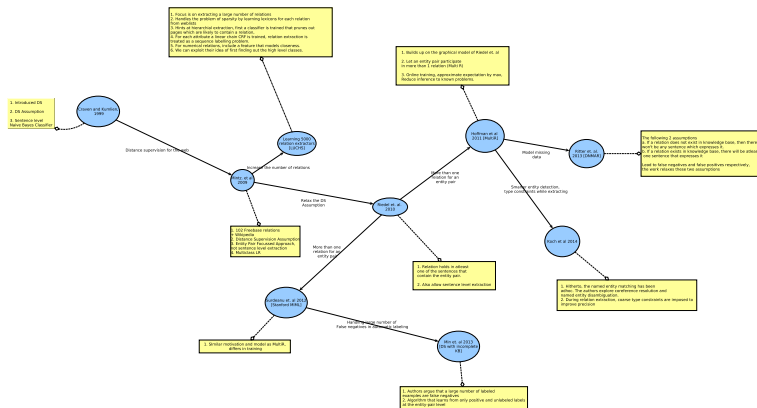


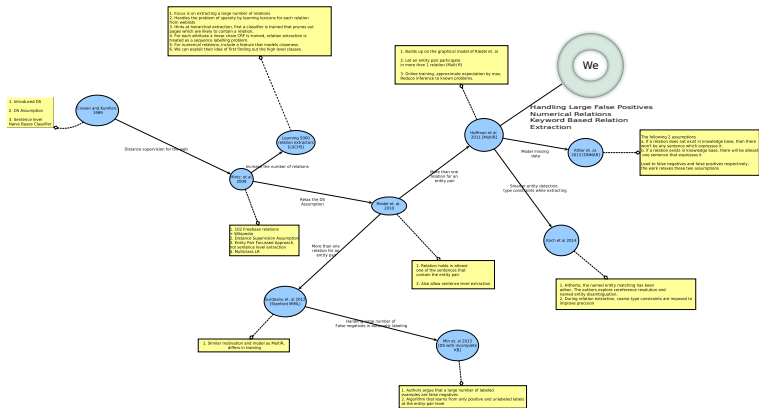
Distant Supervision Techniques

- First paper in 1999, almost every possibility explored



Distant Supervision Techniques

- First paper in 1999, *almost* every possibility explored

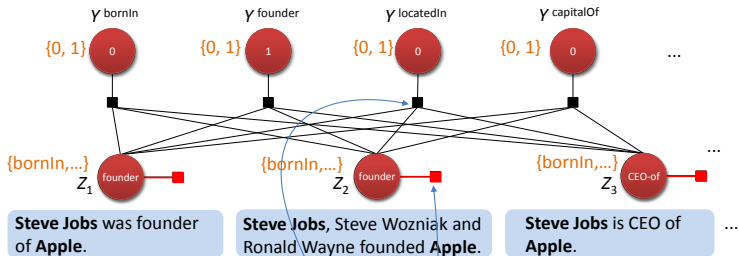


Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Model

Steve Jobs, Apple:



$$p(\mathbf{Y} = \mathbf{y}, \mathbf{Z} = \mathbf{z} | \mathbf{x}; \theta) \stackrel{\text{def}}{=} \frac{1}{Z_{\mathbf{x}}} \prod_r \Phi^{\text{join}}(y^r, \mathbf{z}) \prod_i \Phi^{\text{extract}}(z_i, x_i)$$

$$\Phi^{\text{join}}(y^r, \mathbf{z}) \stackrel{\text{def}}{=} \begin{cases} 1 & \text{if } y^r = \text{true} \wedge \exists i : z_i = r \\ 0 & \text{otherwise} \end{cases}$$

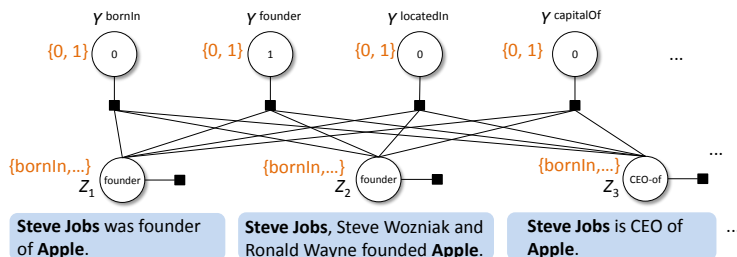
All features at sentence-level

(join factors are deterministic ORs)

Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Model



- Extraction almost entirely driven by sentence-level reasoning
- Tying of facts Y_r and sentence-level extractions Z_i still allows us to model weak supervision for training


Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Inference

Need:


- Most likely sentence labels:



$$\arg \max_{\mathbf{y}, \mathbf{z}} p(\mathbf{y}, \mathbf{z} | \mathbf{x}; \theta)$$

Easy

- Most likely sentence labels *given* facts:



$$\arg \max_{\mathbf{z}} p(\mathbf{z} | \mathbf{x}, \mathbf{y}; \theta)$$

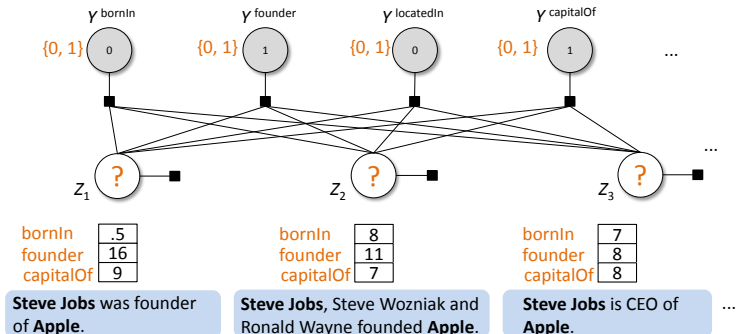
Challenging

Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Inference

- Computing $\arg \max_{\mathbf{z}} p(\mathbf{z}|\mathbf{x}, \mathbf{y}; \theta)$:

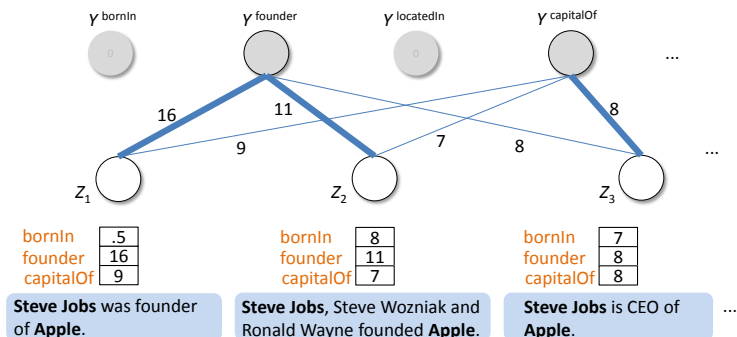


Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Inference

- Variant of the weighted, edge-cover problem:



Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Learning

- Training set $\{(\mathbf{x}_i, \mathbf{y}_i) | i = 1 \dots n\}$, where
 - i corresponds to a particular entity pair
 - \mathbf{x}_i contains all sentences with mentions of pair
 - \mathbf{y}_i bit vector of facts about pair from database
- Maximize Likelihood

$$O(\theta) = \prod_i p(\mathbf{y}_i | \mathbf{x}_i; \theta) = \prod_i \sum_{\mathbf{z}} p(\mathbf{y}_i, \mathbf{z} | \mathbf{x}_i; \theta)$$

Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Learning

- Scalability: Perceptron-style additive updates
- Requires two approximations:
 1. Online learning

For example i (entity pair), define

$$\phi(\mathbf{x}, \mathbf{z}) = \sum_j \phi(x_j, z_j)$$

Use gradient of local log likelihood for example i :

$$\frac{\partial \log O_i(\theta)}{\partial \theta_j} = E_{p(\mathbf{z}|\mathbf{x}_i, \mathbf{y}_i; \theta)}[\phi_j(\mathbf{x}_i, \mathbf{z})] \\ - E_{p(\mathbf{y}, \mathbf{z}|\mathbf{x}_i; \theta)}[\phi_j(\mathbf{x}_i, \mathbf{z})]$$

2. Replace expectations with maximizations

Relation Extraction Using MultiR

From raphaelhoffmann.com/publications

Learning: Hidden-Variable Perceptron

passes over
dataset

for each
entity pair i

most likely
sentence labels
and inferred facts
(ignoring DB facts)

most likely
sentence labels
given DB facts

initialize parameter vector $\Theta \leftarrow 0$

for $t = 1 \dots T$ **do**

for $i = 1 \dots n$ **do**

$(y', z') \leftarrow \arg \max_{y, z} p(y, z | x_i; \theta)$

if $y' \neq y_i$ **then**

$z^* \leftarrow \arg \max_z p(z | x_i, y_i; \theta)$

$\Theta \leftarrow \Theta + \phi(x_i, z^*) - \phi(x_i, z')$

end if

end for

end for

Return Θ